

7.8 Read-Write locks

- ◆ Support one-writer & many readers
- ◆ lockShared() // Reader
- ◆ unlockShared() // Reader
- ◆ lockExclusive() // Writer
- ◆ unlockExclusive() // Writer
- ◆ upgrade() // from shared to exclusive
- ◆ downgrade() //from exclusive to shared

41

Design consideration

- ◆ Avoid needless wake up
- ◆ Reader release
 - ◆ only the last to wake up a single writer
- ◆ Writer release:
 - ◆ prefer writer
 - ◆ Reader starvation
 - ◆ wake up all the readers
 - ◆ Writer starvation
 - ◆ lockShared() block if there is any writer waiting
//new readers

42

Implementation

```
struct rwlock{
int nActive; // number of readers, -1 means a writer is active
int nPendingReads;
int nPendingWrites;
spinlock_t sl;
condition canRead;
condition canWrite;
};
```

43

un/lockShared

```
void lockShared(struct rwlock *r) {
    spin_lock(&r->sl);
    r->nPendingReads++;
    if (r->nPendingWrites>0)
        wait(&r->canRead, &r->sl);
    while(r->nActive <0)
        wait(&r->canRead, &r->sl);
    r->nActive++;
    r->nPendingReads--;
    spin_unlock(&r->sl);
}

void unlockShared(struct rwlock *r) {
    spin_lock(&r->sl);
    r->nActive--;
    if (r->nActive ==0) {
        spin_unlock (&r->sl);
        do_signal(&r->canWrite);
    } else
        spin_unlock(&r->sl);
}
```

44

un/lockExclusive

```

void lockExclusive(struct rwlock *r) {
    spin_lock(&r->sl);
    r->nPendingWrites++;
    while(r->nActive)
        wait(&r->canWrite, &r->sl);
    r->nPendingWrites--;
    r->nActive = 1;
    spin_unlock(&r->sl);
}

void unlockExclusive(struct rwlock *r) {
    boolean_t wakeReaders;
    spin_lock(&r->sl);
    r->nActive = 0;
    wakeReaders = (r->nPendingReads != 0);
    spin_unlock(&r->sl);
    if (wakeReaders) {
        do_broadcast(&r->canRead);
    }
    else
        do_signal(&r->canWrite);
}

```

45

Up/downgrade()

```

void downgrade(struct rwlock *r) {
    boolean_t wakeReaders;
    spin_lock(&r->sl);
    r->nActive = 1;
    wakeReaders =
        (r->nPendingReads != 0);
    spin_unlock(&r->sl);
    if (wakeReaders) {
        do_broadcast(&r->canRead);
    }
}

void upgrade(struct rwlock *r) {
    spin_lock(&r->sl);
    if (r->nActive == 1) r->nActive = -1;
    else{
        r->nPendingWrites++;
        r->nActive--;
        while (r->nActive)
            wait(&r->canWrite, &r->sl);
        r->nPendingWrites--;
        r->nActive = -1;
    }
    spin_unlock(&r->sl);
}

```

46

Using R/W locks

```

rwlock l;

T1() {
    lockShared(&l);
    reading;
    upgrade(&l)
    writing;
    downgrade(&l)
    unlockShared(&l);
}

T2() {
    lockExclusive(&l);
    writing;
    downgrade(&l);
    reading;
    upgrade(&l);
    unlockExclusive(&l);
}

```

47